

archeological site prediction

May 7, 2018

1 Archeological Site Prediction

This notebook predicts the locations of archeological sites in the neolithic era using different types of machine learning models, mostly classification models

1.1 Import data on site, nonsite, and prediction csvs

may need to drop a few columns

```
In [1]: from randomSiteSampling import randomSamplingFromDF
```

```
In [2]: import numpy as np
import pandas as pd
import sklearn as ml
import matplotlib
%matplotlib inline
```

```
import pprint
pp = pprint.PrettyPrinter(indent=4).pprint
```

```
"""
```

```
Importing data
```

```
"""
```

```
raw_site_file = 'common_sites_V3.csv'
raw_pred_file = 'ireland_sites_V3.csv'
```

```
raw_site_df = pd.read_csv(raw_site_file, low_memory=False, error_bad_lines=False, encoding='utf-8')
raw_pred_df = pd.read_csv(raw_pred_file, low_memory=False, error_bad_lines=False, encoding='utf-8')
```

```
site_df = raw_site_df.copy()
nonsite_df, pred_df = randomSamplingFromDF(raw_site_df, raw_pred_df)
```

```
ireland_site_df = site_df.where(site_df['country'] == 'Ireland')
```

```
"""
```

```
Making sure all dataframes have the same columns
```

```

We do this by finding all the columns which are
not shared between all dataframes. We then keep
only the columns common to all the dataframes
"""
site_attr      = set(site_df)
nonsite_attr   = set(nonsite_df)
pred_attr      = set(pred_df)

common_attr    = site_attr & nonsite_attr & pred_attr
uncommon_attr  = (site_attr | nonsite_attr | pred_attr) - common_attr

# create list of columns to drop
drop_site      = site_attr & uncommon_attr
drop_nonsite   = nonsite_attr & uncommon_attr
drop_pred      = pred_attr & uncommon_attr

# drop columns
site_df        = site_df.drop(columns=drop_site)
nonsite_df     = nonsite_df.drop(columns=drop_nonsite)
pred_df        = pred_df.drop(columns=drop_pred)

"""
Take care of na values by propogating values forward then backward
"""
site_df        = site_df.fillna (method='ffill').fillna(method='bfill')
nonsite_df     = nonsite_df.fillna(method='ffill').fillna(method='bfill')
pred_df        = pred_df.fillna (method='ffill').fillna(method='bfill')

```

2 Combining our site and nonsite data

```

In [3]: site_df['site'] = 1
        nonsite_df['site'] = 0
        data = pd.concat([site_df, nonsite_df])

        target = data['site']
        data = data.drop(columns=['site'])

```

3 Splitting the Data into Training and Test Sets

```

In [4]: from sklearn.model_selection import train_test_split

```

```

train_size=0.7

```

```

train_data, test_data, train_target, test_target = train_test_split(data, target, train_

```

```

/Users/WillC/miniconda2/envs/ML/lib/python2.7/site-packages/sklearn/model_selection/_split.py:20
FutureWarning)

```

4 Train using Different Models

1. SVC Model (Linear kernel)

2. K-Neighbors Classifier

4. Decision Forest Classifier

4. Random Forest Classifier

5. Gradient Boosting Classifier

6. Ada Boosting Classifier

4.0.1 The 'predictions' dataframe will be used to store the results of the following training models

```
In [5]: predictions = pred_df.copy(deep=True)
```

```
In [6]: from plotPredicted import plotPredictedFromDF
```

4.1 1. Linear SVC Model

Training the Linear SVC Model

```
In [7]: from sklearn import svm
```

```
SVCModel = svm.SVC(kernel='linear', C=.5)
SVCModel.fit(train_data, train_target)
```

```
Out[7]: SVC(C=0.5, cache_size=200, class_weight=None, coef0=0.0,
          decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
          max_iter=-1, probability=False, random_state=None, shrinking=True,
          tol=0.001, verbose=False)
```

Testing the Linear SVC Model

```
In [8]: from sklearn import metrics
```

```
model = SVCModel
"""
Fit the training data
and observe the metrics
"""
preds = model.predict(train_data)
```

```

targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy  : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall    : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1        : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy  : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall    : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1        : %.03f%s" % (metrics.f1_score(*args)*100, "%")

```

```

Training Data
accuracy  : 84.713%
precision : 93.526%
recall    : 75.046%
f1        : 83.273%

```

```

Test Data
accuracy  : 83.991%
precision : 90.859%
recall    : 74.376%
f1        : 81.796%

```

Linear SVC Prediction

```

In [43]: model_name = 'LinearSVCPrediction'
         predictions[model_name] = SVCModel.predict(pred_df)

         site_exists = predictions[predictions[model_name] == 1]
         site_nexists = predictions[predictions[model_name] != 1]
         num_site_exists = site_exists[model_name].count()
         num_site_nexists = site_nexists[model_name].count()

         print 'Sites exists          : %d ' % (num_site_exists)
         print 'Sites does not exists: %d ' % (num_site_nexists)

```

```

print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
(100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

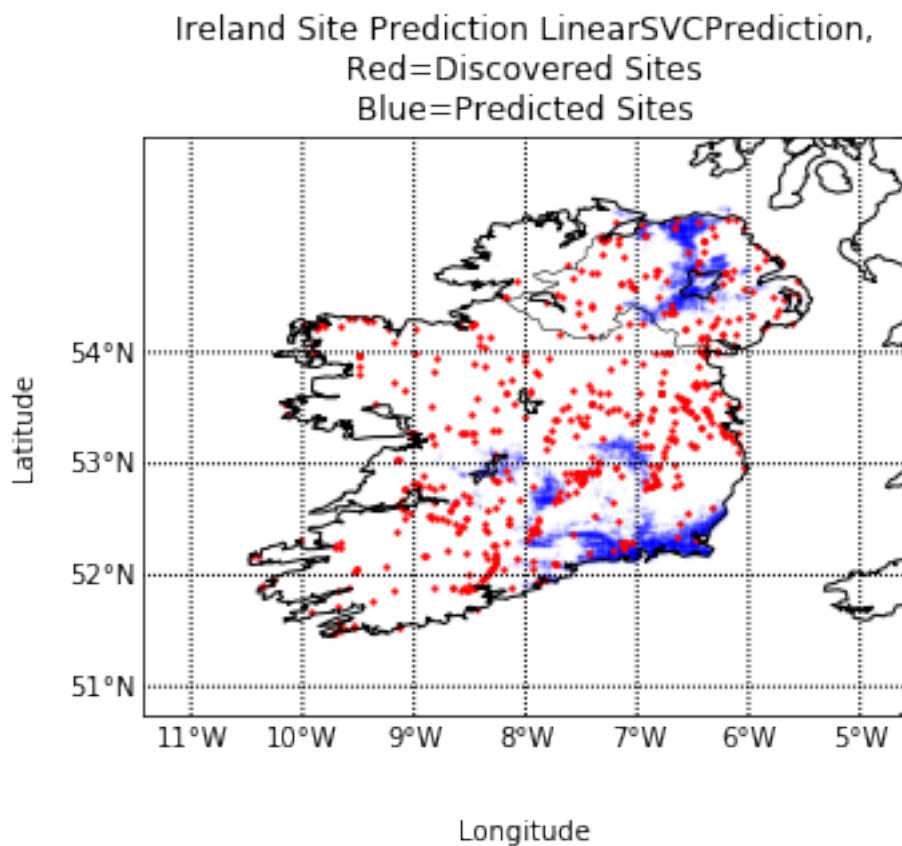
# plot the map
plotPredictedFromDF(raw_site_df,
                    predictions,
                    model_name=model_name,
                    country='Ireland',
                    resolution='i',
                    alpha_predicted=.01)

```

Sites exists : 40491

Sites does not exists: 618446

Percentage of sites predicted having an archeological site: 6.145%



4.2 2. K-Neighbors Classifier

Training the K-Neighbors Classifier Model

```
In [21]: from sklearn import neighbors

n_neighbors=20
KNeighborsModel = neighbors.KNeighborsClassifier(n_neighbors=n_neighbors)
KNeighborsModel.fit(train_data, train_target)

Out[21]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=1, n_neighbors=20, p=2,
                             weights='uniform')
```

Testing the K-Neighbors Classifier Model

```
In [22]: from sklearn import metrics

model = KNeighborsModel
"""
Fit the training data
and observe the metrics
"""

preds = model.predict(train_data)
targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")
```

```
Training Data
accuracy : 84.760%
precision : 97.243%
recall : 71.985%
f1 : 82.729%
```

Test Data
accuracy : 84.649%
precision : 96.308%
recall : 70.975%
f1 : 81.723%

K-Neighbors Prediction

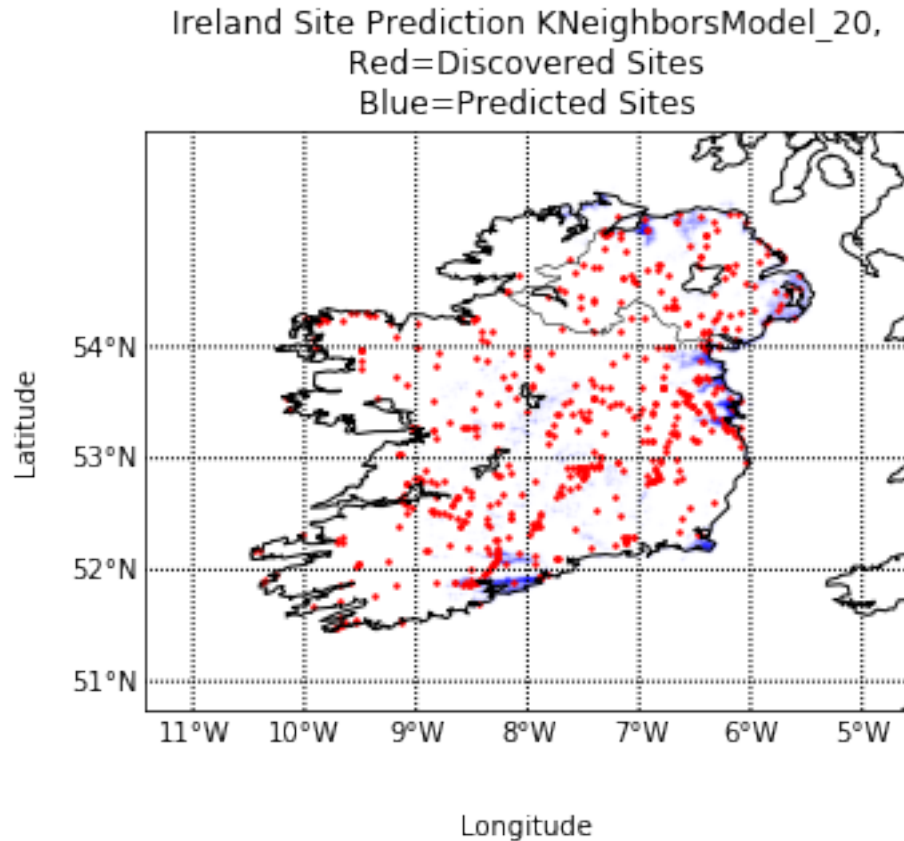
```
In [42]: model_name = 'KNeighborsModel_%d' % n_neighbors
         predictions[model_name] = KNeighborsModel.predict(pred_df)

         site_exists      = predictions[predictions[model_name] == 1]
         site_nexists     = predictions[predictions[model_name] != 1]
         num_site_exists  = site_exists[model_name].count()
         num_site_nexists = site_nexists[model_name].count()

         print 'Sites exists          : %d ' % (num_site_exists)
         print 'Sites does not exists: %d ' % (num_site_nexists)
         print 'Percentage of sites predicted having an archeological site'
         print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
         (100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

         # plot the map
         plotPredictedFromDF(raw_site_df,
                             predictions,
                             model_name=model_name,
                             country='Ireland',
                             resolution='i',
                             alpha_predicted=.01)

Sites exists          : 15240
Sites does not exists: 643697
Percentage of sites predicted having an archeological site
Percentage of sites predicted having an archeological site: 2.313%
```



4.3 3. Decision Tree Classification

In [25]: `from sklearn import tree`

```
DecisionTree = tree.DecisionTreeClassifier()
DecisionTree.fit(train_data, train_target)
```

Out[25]: `DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, presort=False, random_state=None, splitter='best')`

Testing the Decision Tree Classifier Model

In [26]: `from sklearn import metrics`

```
model = DecisionTree
"""
```



```

Fit the training data
and observe the metrics
"""

preds = model.predict(train_data)
targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

```

```

Training Data
accuracy : 100.000%
precision : 100.000%
recall : 100.000%
f1 : 100.000%

```

```

Test Data
accuracy : 93.311%
precision : 92.411%
recall : 93.878%
f1 : 93.138%

```

Decision Tree Prediction

```

In [38]: model_name = 'DecisionTree'
         predictions[model_name] = DecisionTree.predict(pred_df)

         site_exists      = predictions[predictions[model_name] == 1]
         site_nexists     = predictions[predictions[model_name] != 1]
         num_site_exists  = site_exists[model_name].count()

```

```

num_site_nexists = site_nexists[model_name].count()

print 'Sites exists          : %d ' % (num_site_exists)
print 'Sites does not exists: %d ' % (num_site_nexists)
print 'Percentage of sites predicted having an archeological site'
print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
(100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

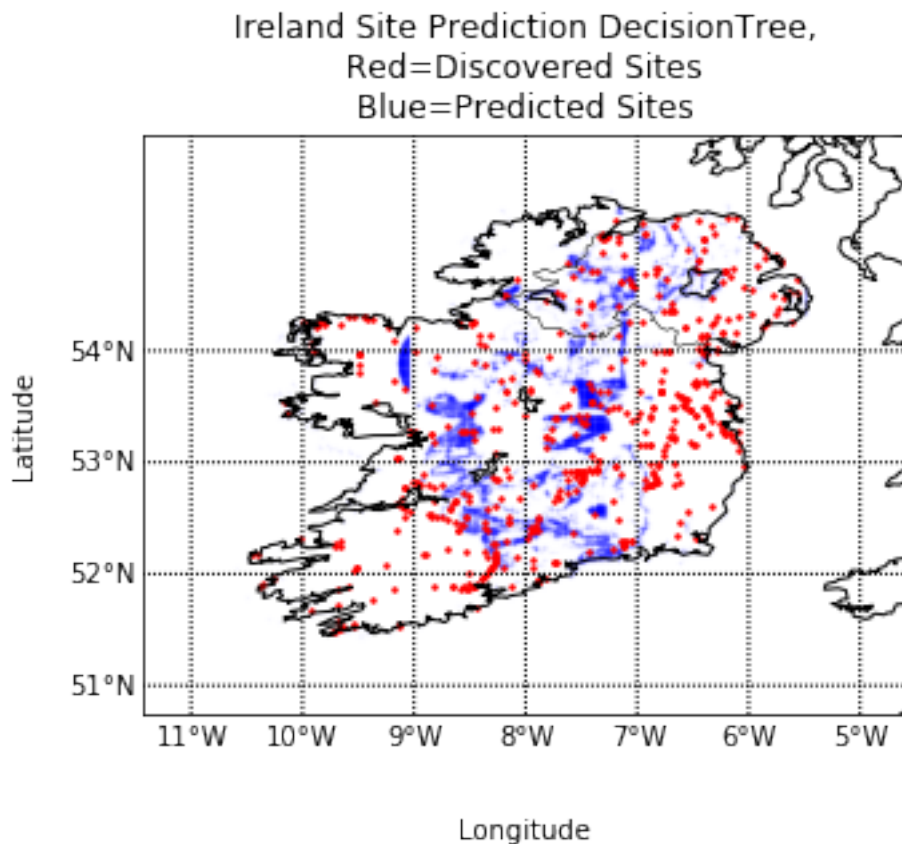
# plot the map
plotPredictedFromDF(raw_site_df,
                    predictions,
                    model_name=model_name,
                    country='Ireland',
                    resolution='i',
                    alpha_predicted=.01)

```

```

Sites exists          : 51006
Sites does not exists: 607931
Percentage of sites predicted having an archeological site
Percentage of sites predicted having an archeological site: 7.741%

```



4.4 4. Random Forest Classification

```
In [28]: from sklearn import ensemble
```

```
n_estimators=10
RandomForest = ensemble.RandomForestClassifier(n_estimators=10)
RandomForest.fit(train_data, train_target)
```

```
Out[28]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

Testing the Random Forest Classifier Model

```
In [29]: from sklearn import metrics
```

```
model = RandomForest
"""
Fit the training data
and observe the metrics
"""

preds = model.predict(train_data)
targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")
```

```
Training Data
accuracy : 99.577%
```

```
precision : 100.000%
recall    : 99.165%
f1        : 99.581%
```

Test Data

```
accuracy : 93.969%
precision : 98.010%
recall    : 89.342%
f1        : 93.476%
```

Random Forest Prediction

```
In [39]: model_name = 'RandomForest_%d' % (n_estimators)
        predictions[model_name] = RandomForest.predict(pred_df)

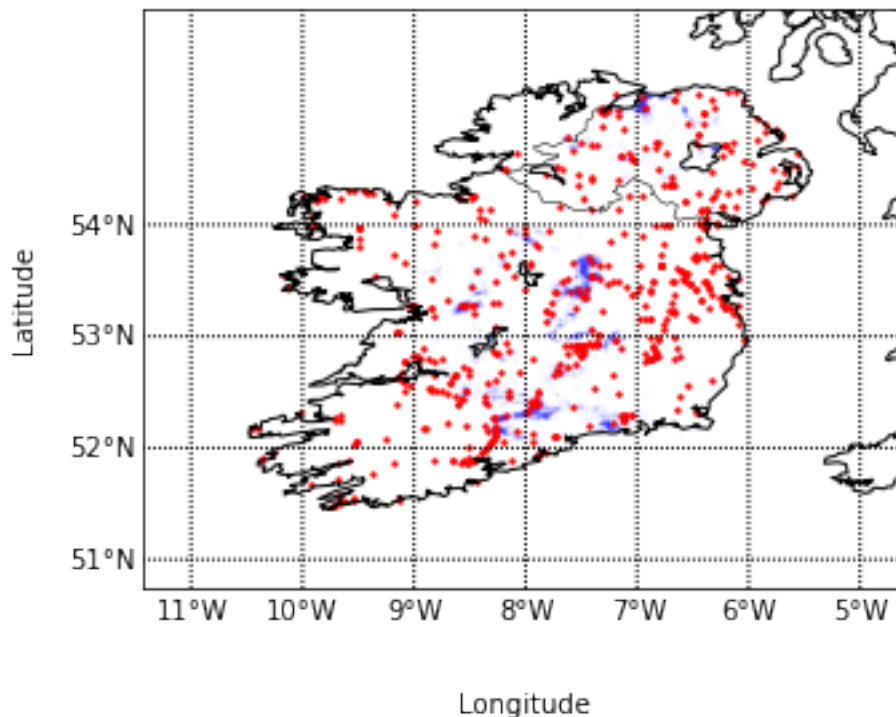
        site_exists      = predictions[predictions[model_name] == 1]
        site_nexists     = predictions[predictions[model_name] != 1]
        num_site_exists  = site_exists[model_name].count()
        num_site_nexists = site_nexists[model_name].count()

        print 'Sites exists          : %d ' % (num_site_exists)
        print 'Sites does not exists: %d ' % (num_site_nexists)
        print 'Percentage of sites predicted having an archeological site'
        print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
        (100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

        # plot the map
        plotPredictedFromDF(raw_site_df,
                           predictions,
                           model_name=model_name,
                           country='Ireland',
                           resolution='i',
                           alpha_predicted=.01)
```

```
Sites exists          : 11342
Sites does not exists: 647595
Percentage of sites predicted having an archeological site
Percentage of sites predicted having an archeological site: 1.721%
```

Ireland Site Prediction RandomForest_10,
Red=Discovered Sites
Blue=Predicted Sites



4.5 5. Gradient Boosting Classification

In [31]: `from sklearn import ensemble`

```
GradientBoosting = ensemble.GradientBoostingClassifier()
GradientBoosting.fit(train_data, train_target)
```

Out[31]: `GradientBoostingClassifier(criterion='friedman_mse', init=None, learning_rate=0.1, loss='deviance', max_depth=3, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, n_estimators=100, presort='auto', random_state=None, subsample=1.0, verbose=0, warm_start=False)`

Testing the Gradient Boosting Classifier

In [32]: `from sklearn import metrics`

```

model = GradientBoosting
"""
Fit the training data
and observe the metrics
"""

preds = model.predict(train_data)
targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1 : %.03f%s" % (metrics.f1_score(*args)*100, "%")

```

```

Training Data
accuracy : 97.131%
precision : 100.000%
recall : 94.341%
f1 : 97.088%

```

```

Test Data
accuracy : 95.504%
precision : 98.780%
recall : 91.837%
f1 : 95.182%

```

Gradient Boosting Classifier Prediction

```

In [40]: model_name = 'GradientBoosting'
         predictions[model_name] = GradientBoosting.predict(pred_df)

         site_exists      = predictions[predictions[model_name] == 1]

```

```

site_nexists      = predictions[predictions[model_name] != 1]
num_site_exists   = site_exists[model_name].count()
num_site_nexists  = site_nexists[model_name].count()

print 'Sites exists          : %d ' % (num_site_exists)
print 'Sites does not exists: %d ' % (num_site_nexists)
print 'Percentage of sites predicted having an archeological site:'
print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
(100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

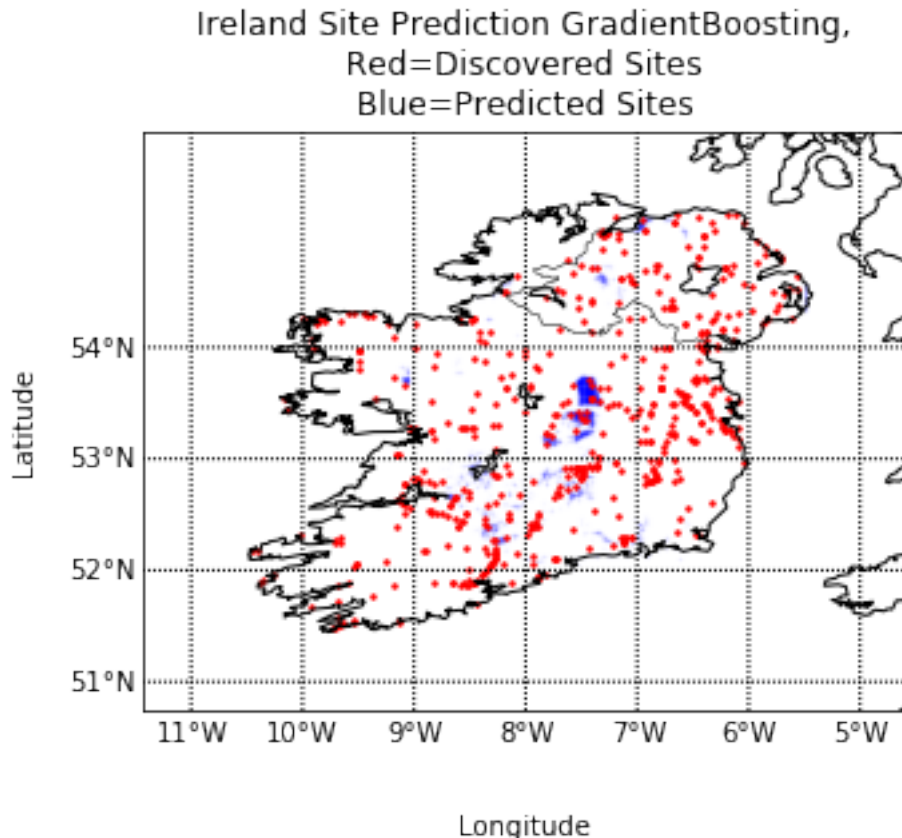
# plot the map
plotPredictedFromDF(raw_site_df,
                    predictions,
                    model_name=model_name,
                    country='Ireland',
                    resolution='i',
                    alpha_predicted=.01)

```

```

Sites exists          : 8044
Sites does not exists: 650893
Percentage of sites predicted having an archeological site:
Percentage of sites predicted having an archeological site: 1.221%

```



4.6 6. Ada Boosting Classification

```
In [35]: from sklearn import ensemble
```

```
AdaBoosting = ensemble.AdaBoostClassifier()
AdaBoosting.fit(train_data, train_target)
```

```
Out[35]: AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None,
                             learning_rate=1.0, n_estimators=50, random_state=None)
```

Testing the Ada Boosting Classification

```
In [36]: from sklearn import metrics
```

```
model = AdaBoosting
"""
Fit the training data
and observe the metrics
"""

preds = model.predict(train_data)
targs = train_target
args = (targs, preds)
print "Training Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall    : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1        : %.03f%s" % (metrics.f1_score(*args)*100, "%")

"""
Fit the test data
and observe the metrics
"""

preds = model.predict(test_data)
targs = test_target
args = (targs, preds)
print "\n"
print "Test Data"
print "accuracy : %.03f%s" % (metrics.accuracy_score(*args)*100, "%")
print "precision : %.03f%s" % (metrics.precision_score(*args)*100, "%")
print "recall    : %.03f%s" % (metrics.recall_score(*args)*100, "%")
print "f1        : %.03f%s" % (metrics.f1_score(*args)*100, "%")
```

```
Training Data
accuracy : 95.296%
precision : 98.225%
```



```
recall    : 92.393%
f1        : 95.220%
```

Test Data

```
accuracy  : 91.776%
precision : 92.558%
recall    : 90.249%
f1        : 91.389%
```

Ada Boosting Prediction

```
In [41]: model_name = 'AdaBoosting'
        predictions[model_name] = AdaBoosting.predict(pred_df)

        site_exists      = predictions[predictions[model_name] == 1]
        site_nexists     = predictions[predictions[model_name] != 1]
        num_site_exists  = site_exists[model_name].count()
        num_site_nexists = site_nexists[model_name].count()

        print 'Sites exists          : %d ' % (num_site_exists)
        print 'Sites does not exists: %d ' % (num_site_nexists)
        print 'Percentage of sites predicted having an archeological site: %.3f%s' % \
              (100.0*float(num_site_exists)/float(num_site_exists+num_site_nexists), '%')

        # plot the map
        plotPredictedFromDF(raw_site_df,
                             predictions,
                             model_name=model_name,
                             country='Ireland',
                             resolution='i',
                             alpha_predicted=.01)
```

```
Sites exists          : 34408
Sites does not exists: 624529
Percentage of sites predicted having an archeological site: 5.222%
```

Ireland Site Prediction AdaBoosting,
Red=Discovered Sites
Blue=Predicted Sites

